

Micro-benchmarks: A comparison of C++, Rust, and Zig

Fucă Răzvan-Ionuț
Technical University of Cluj-Napoca

Contents

1	Introduction	3
1.1	Context	3
1.2	Language selection	3
1.3	Objectives	3
1.4	Hypothesis	3
2	Bibliographic Study	4
2.1	Introduction to the languages	4
2.2	Visualization tools	4
3	Analysis & Design	5
3.1	Memory benchmarks	5
3.2	Threading benchmarks	5
4	Implementation	7
4.1	Memory tests, dynamic data	7
4.2	Memory tests, static data	7
4.3	Threading tests, thread creation	7
4.4	Threading tests, thread context switch	7
4.5	Threading tests, thread migration	7
4.6	Benchmarking script	8
5	Testing and validation	9
5.1	Interpretation of results	9
6	Conclusion	11
7	Possible Extensions	11
8	Bibliography	12

1 Introduction

1.1 Context

The purpose of this project is to conduct a comprehensive benchmark of the execution time of various processes in three distinct programming languages: C++, Rust, and Zig. These languages have been chosen due to their classification as "systems programming languages," which means they offer direct access to hardware while also providing a robust, OS-agnostic API. Notably, they all support multi-threading, making them suitable for low-level and high-performance tasks. This project will delve into the nuances of these languages, shedding light on their capabilities and performance characteristics.

1.2 Language selection

The languages chosen for this project are:

- C++ - A widely used systems programming language that offers extensive control over memory management and hardware interactions. C++ has a long history and a rich ecosystem of libraries and tools.
- Rust - A modern systems programming language known for its strong focus on memory safety without sacrificing performance. Rust's unique ownership system and borrowing rules set it apart from traditional languages.
- Zig - An up-and-coming systems programming language that combines a C-like syntax with modern features and a focus on compile-time safety. Zig's custom allocators and comptime code are key features that distinguish it.

1.3 Objectives

The primary objectives of this project are:

- Develop sample programs in each of the selected languages that are specifically designed to test the performance of critical processes, including:
 - Memory allocation
 - Static memory access
 - Dynamic memory access
 - Thread creation
 - Thread context switch
 - Thread migration
- Execute the sample programs in order to collect and interpret the data.

1.4 Hypothesis

The hypothesis for this project is that the results of the benchmarks will demonstrate close performance characteristics among the selected languages. This expectation is based on the fact that all three languages ultimately compile down to assembly instructions, but differences will arise from:

- Compiler Optimizations: Each language's compiler applies its own set of optimizations, affecting execution speed.
- Programming Paradigms: The choice of language influences the programming paradigms and patterns used, which can impact performance.
- Memory Management Strategies: The languages vary in their approach to memory management, such as C++ relying on manual memory management, Rust utilizing its borrow checker, and Zig offering custom allocators.

2 Bibliographic Study

2.1 Introduction to the languages

The programming languages selected for this micro-benchmarking study are C++, Rust, and Zig. These languages have the ability to provide low-level control over hardware resources and memory management. In this section, we provide an overview of these languages and their relevance to our benchmarking project.

- C++ is a well-established systems programming language known for its versatility and extensive control over memory management and hardware interactions. It has a rich history, starting as an extension of C in 1979.
- Rust is a modern systems programming language that has garnered attention for its emphasis on memory safety without compromising performance. It distinguishes itself with its ownership system and borrowing rules, forcing the programmer to adopt the RAII (resource acquisition is initialization) programming pattern, eliminating common sources of memory-related bugs.
- Zig is an emerging systems programming language that combines a C-like syntax with modern features. It places a strong focus on compile-time safety, which aligns well with our objective to assess language-level safety and performance. Notably, Zig offers features like custom allocators and compile-time code that set it apart from traditional systems programming languages.

2.2 Visualization tools

To effectively present and analyze the data collected from our benchmarking programs, we will utilize Matplotlib, a widely-used Python library for creating high-quality data visualizations. Matplotlib provides a versatile and user-friendly platform for generating graphs and charts, making it an ideal choice for representing the performance metrics obtained during our micro-benchmarking study. Matplotlib has also been chosen because of its ease of processing external data such as CSV files.

3 Analysis & Design

The programs will be CLI apps obeying *NIX philosophy that will perform a series of algorithms designed to benchmark the aforementioned processes and will dump the results to a file in a plot-friendly format (CSV) that will later be processed by a script that will prune and interpret the data. Figure 1 showcases the general flow of data in our project.

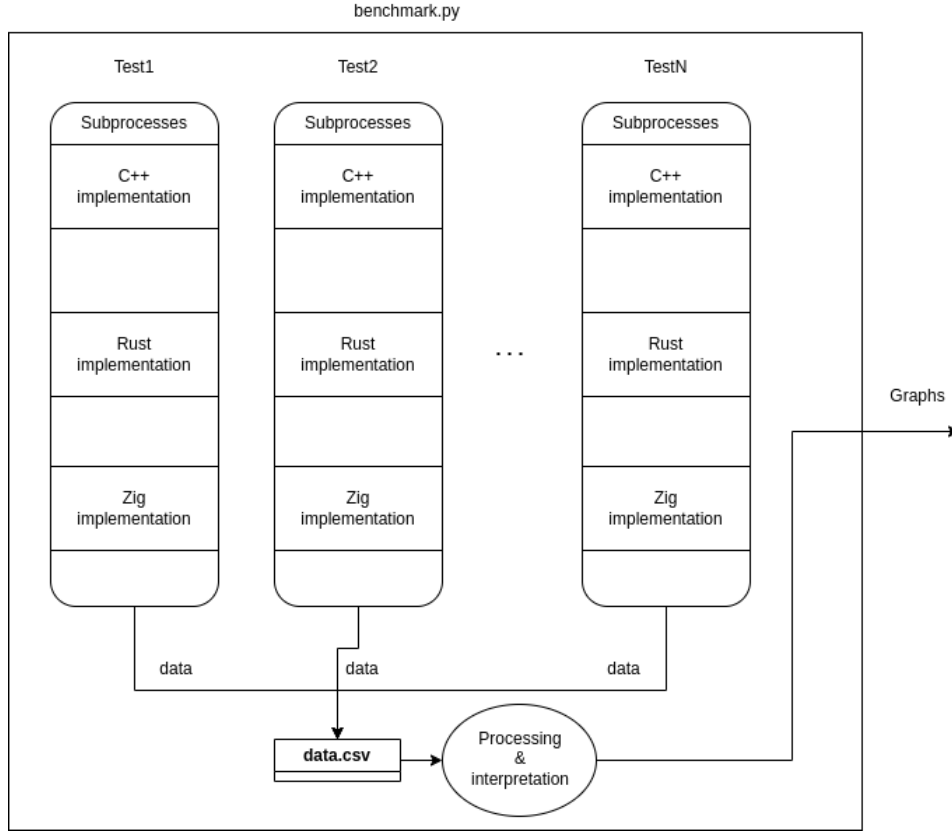


Figure 1: General flow of data in the project.

The type of benchmarks that will be performed can be grouped quite easily into two main categories:

- Memory benchmarks
- Threading benchmarks

3.1 Memory benchmarks

For the memory benchmarks, the programs will perform a series elementary operations on increasingly larger common data structures. A complete list of actions to be performed that I believe will encompass most common applications is:

- Allocating dynamic arrays
- Creating linked lists (or similar dynamic data structures)
- Performing operations such as sorting or traversals on the created data structures

3.2 Threading benchmarks

For the threading benchmarks, the measuring process will be a little more difficult as there is no explicit way of directly accomplishing our desired results; however, there are workarounds that will force the OS to perform operations such as context switching and migration. A list of actions that I believe will adequately measure the desired processes consists of:

- Calling an empty function in a synchronous and asynchronous way and computing the difference between the two to gather quantitative data regarding thread creation
- Writing and reading bytes of data from a pipe to force two thread context switches
- Setting thread CPU affinity using functions like `pthread_set_affinity_np` or similar depending on the platform to force thread migration

All of these benchmarks will be time-based, using each of the language's built-in time measuring method that will yeild results in nanoseconds:

- C++ `std::chrono::high_resolution_clock`
- Zig's `std.time.Timer`
- Rust's `std::time::Instant`

4 Implementation

4.1 Memory tests, dynamic data

The first step of running our benchmarks is implementing the necessary data structures we need to test. The most notable of the data structures is a generic singly linked list that will showcase each language's lower-level features (manual memory allocation, the use of pointers and nullable/optional data types, etc.), as well as higher-level features (OOP, static and dynamic dispatch, or compile-time code), and the programming paradigm they encourage.

The implementation details of each are described below:

- **C++** – Singly linked list using templates
- **Zig** – Compile-time function that returns a new type
- **Rust** – Generic list that uses the 'Box' smart pointer for dynamic memory allocation

Another test relies on the dynamic allocation of large arrays on the heap using each of the languages specific constructs

- **C++** – Using the new operator
- **Zig** – Using different allocators provided by the standard library
- **Rust** – Using `Vec::with_capacity` as it is the closest analogue to allocating a dynamic array of a known size in rust

4.2 Memory tests, static data

Each program will allocate an array on the stack of size 100000 and then traverse it, setting the value of each entry in the array to their index.

4.3 Threading tests, thread creation

Each program will run a function that returns as soon as possible. At first by launching it in a thread, and secondly by running it directly. We will measure the time difference between these two scenarios in order to quantify thread creation.

The rust implementation also provides the benchmarking of a 'Green thread' using the tokio crate This will be accomplished using:

- **C++** – `std::thread`
- **Zig** – `std.Thread`
- **Rust** – `std::thread`
- **Rust green threads** – `tokio::task::spawn`

4.4 Threading tests, thread context switch

The programs will create a child thread and ping_pong a single byte between parent and child. This will force two thread context switches per read/write

4.5 Threading tests, thread migration

The programs will leverage `pthread_set_affinity_np` in order to force the tested thread to run on a single particular processor. We will force the thread to migrate from processor one to processor two using the linux system call `sched_yield`.

4.6 Benchmarking script

The script that will run and aggregate the data generated by the programs is `benchmark.py`. This benchmarking script relies on a couple of files

- **config.yaml** Where the benchmarks parameters will be defined. Ex (Size of array, number of runs, weather or not to recompile the programs) etc.
- **data.csv** where all the programs will dump their results and will act as the data aggregator for the whole project

The script also provides a GUI in if the user prefers to use it instead of a config file.

This script is responsible for compiling the tests of every language (if so desired), running them and interpreting the results by creating graphs from the provided data

All the results of the tests are saved into the `charts/` folder and use the current unix timestamp to differentiate between them

5 Testing and validation

5.1 Interpretation of results

As anticipated, the majority of the results exhibit comparability, with variations across each run attributed to statistical differences. Noteworthy disparities emerge in tests that involve dynamic memory allocation, showcasing Zig’s precise control over memory through allocators. Additionally, significant distinctions appear in thread creation tests, underscoring the unmatched efficiency of green threads

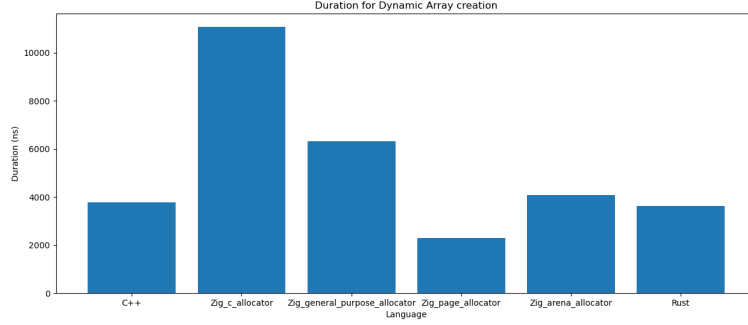


Figure 2: Dynamic array creation

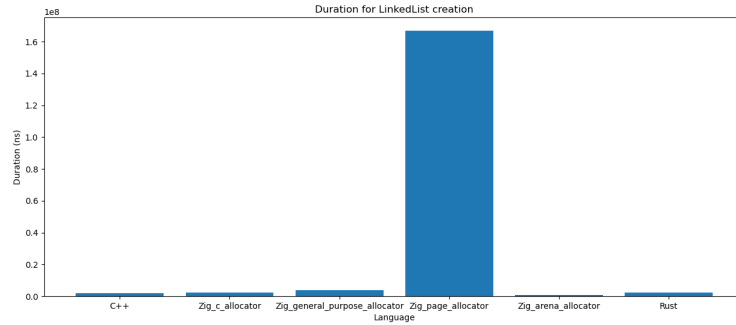


Figure 3: Linked List creation

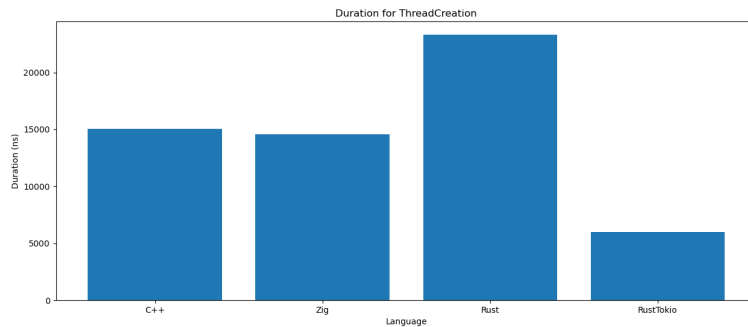


Figure 4: Thread Creation

Upon analyzing the aforementioned graphs, it becomes evident that no one-size-fits-all tool exists, even for seemingly elementary operations like memory allocation.

The Operating System’s page allocator excels in allocating large blocks of contiguous data but exhibits suboptimal performance for frequent small allocations, such as those required by dynamic data structures like linked lists.

On the other hand, a compelling choice for a general-purpose allocator is the arena allocator, which acquires a substantial chunk of memory from the OS (the arena). Subsequent calls then utilize this arena as a "pool" of available memory. This approach proves to be highly efficient for both large and small allocations. Furthermore, it offers the advantage of deallocating the entire arena when it is no longer needed, mitigating a spectrum of potential memory-related issues that may arise from manual memory management (e.g., double-free, call-after-free).

Green threads, often referred to as lightweight or user-level threads, are threads managed entirely by a user-space runtime or library rather than the operating system. They provide a higher level of abstraction compared to traditional operating system threads. The "lightweight" aspect can be clearly seen in the above graph as Rust's dominant green-thread library, Tokio, is able to consistently provide the fastest startup times when it comes to creating a new thread

Of course, different computers may get different results. The computer that the tests were ran on has the following config (extracted from neofetch):

```
OS: Ubuntu 23.10 x86_64
Host: 81Y6 Lenovo Legion 5 15IMH05H
Kernel: 6.5.0-14-generic
CPU: Intel i5-10300H (8) @ 4.500GHz
GPU: NVIDIA GeForce RTX 2060 Mobile
Memory: 7576MiB / 15891MiB
```

6 Conclusion

The goal of the project was to benchmark the execution times of several processes in different programming languages.

As such, this project uses the following programming languages

1. *Python* as a general purpose scripting language that compiles and interprets data while also providing several user-friendly UI options
2. *Rust* Used for benchmarks, sets itself apart with ownership system and functional programming inspired uses
3. *Zig* Used for benchmarks, sets itself apart by the minute level of control over all aspects of the code and by its ability to run code at compile-time
4. *C++* One of, if not the most used general purpose programming language used to date.

7 Possible Extensions

The project presents ample opportunities for expansion, and the following features represent just a glimpse into the potential avenues for further development:

1. Diversified Programming Language Support:

- **Interpreted and Compiled Languages:** Expand the benchmark suite to include a broader array of programming languages, covering both interpreted (e.g., Python, Ruby) and compiled languages (e.g., Go, C#).
- **Strongly Typed and Weakly Typed Languages:** Introduce benchmarks for languages with different type systems, exploring the performance characteristics of strongly typed (e.g., Java, C#) and weakly typed languages (e.g., JavaScript, Python, Ruby).
- **Paradigms:** Extend support for various programming paradigms, encompassing object-oriented (e.g., Java, C++), procedural (e.g., C), and functional languages (e.g., Haskell, Scala).
- **Runtimes** Extend support for languages that utilize shared runtimes (JVM - Java, Kotlin, Scala, Groovy), .NET (C#, F#), ErlangVM (Erlang, Elixir) etc.

2. Diversified Operation Testing:

- **Disk I/O Performance:** Incorporate benchmarks that assess the efficiency of disk input/output operations, simulating real-world scenarios involving file handling and storage.
- **Networking Performance:**
 - **Throughput and Latency:** Expand testing to include benchmarks for networking throughput and latency, crucial metrics for applications involving communication over networks.
 - **Bandwidth Utilization:** Evaluate the efficiency of programming languages in utilizing network bandwidth, especially relevant for data-intensive applications.

3. Exploration of Higher-Level Features:

- **Generics:** Extend the benchmark suite to evaluate the performance of programming languages in handling generic programming constructs.
- **Polymorphism and Dynamic Dispatch:**
 - **Dynamic Dispatch:** Evaluate the performance implications of dynamic dispatch in object-oriented languages, shedding light on the efficiency of runtime method resolution and their underlying data structures (vTables, fat pointers etc)

These potential extensions not only broaden the scope of the project but also provide valuable insights into the performance characteristics of programming languages across a diverse set of criteria. Continuous exploration and expansion will enhance the versatility of the benchmark suite, making it a comprehensive tool for evaluating and comparing languages in various application scenarios.

8 Bibliography

1. *C++ language reference*. [Online]. Available: <https://en.cppreference.com/w/cpp/language>
2. *Rust language reference*. [Online]. Available: <https://doc.rust-lang.org/stable/reference/>
3. *Zig language reference*. [Online]. Available: <https://ziglang.org/documentation/master/>
4. *Matplotlib documentation*. [Online]. Available: <https://matplotlib.org/stable/index.html>
5. *Rust by example: RAII*. [Online]. Available: <https://doc.rust-lang.org/rust-by-example/scope/raii.html>